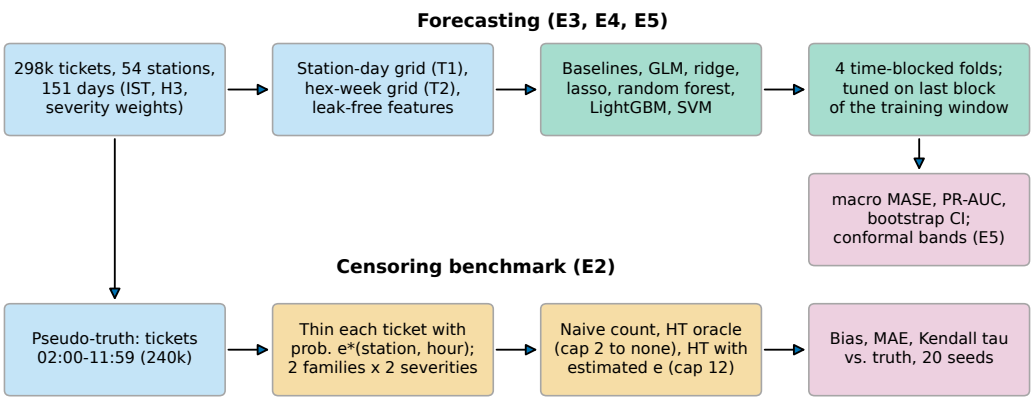
Forecasting (E3, E4, E5)
298k tickets, 54 stations, 151 days (IST, H3, severity weights)
Station-day grid (T1), hex-week grid (T2), leak-free features
Baselines, GLM, ridge, lasso, random forest, LightGBM, SVM
4 time-blocked folds; tuned on last block of the training window
macro MASE, PR-AUC, bootstrap CI; conformal bands (E5)
Censoring benchmark (E2)
Pseudo-truth: tickets 02:00-11:59 (240k)
Thin each ticket with prob. e*(station, hour); 2 families x 2 severities
Naive count, HT oracle (cap 2 to none), HT with estimated e (cap 12)
Bias, MAE, Kendall tau vs. truth, 20 seeds